

Week 6 Assignment

Problem 12.1a Given the classification imbalance in hepatic injury status, describe how you would create a training and testing set

- I would use a stratified random sampling technique because that would be more representative of the proportions in the overall response than to use a simple random sampling technique.

Problem 12.1b Which classification statistic would you choose to optimize for this exercise and why?

- I believe Kappa is a good statistic to use to judge the models here. It is more informative than simply using the accuracy rate, as the accuracy rate simply looks at all classes as equal. The highest Kappa value will be considered the best model. Also since there are more than two outcomes, ROC becomes a little less interpretable, which is why we may choose not to use this metric.

Problem 12.1c Pre-process the data, split the data into a training and a testing set, and build models described in this chapter for the biological predictors. Using each model to predict on the testing set, which model has the best predictive ability for the biological predictors and what is the optimal performance?

- Linear Discriminant Analysis

Linear Discriminant Analysis

225 samples

91 predictor

3 classes: 'Mild', 'None', 'Severe'

Pre-processing: centered (91), scaled (91)

Resampling: Repeated Train/Test Splits Estimated (10 reps, 75%)

Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...

Resampling results:

Accuracy	Kappa
0.4446429	0.08415325

Confusion Matrix and Statistics

Reference			
Prediction	Mild	None	Severe
Mild	16	9	3
None	11	10	1
Severe	2	2	2

Overall Statistics

Accuracy : 0.5
95% CI : (0.3634, 0.6366)
No Information Rate : 0.5179
P-Value [Acc > NIR] : 0.6562

Kappa : 0.1413

McNemar's Test P-Value : 0.8653

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.5517	0.4762	0.33333
Specificity	0.5556	0.6571	0.92000
Pos Pred Value	0.5714	0.4545	0.33333
Neg Pred Value	0.5357	0.6765	0.92000
Prevalence	0.5179	0.3750	0.10714
Detection Rate	0.2857	0.1786	0.03571
Detection Prevalence	0.5000	0.3929	0.10714
Balanced Accuracy	0.5536	0.5667	0.62667

-

- Partial Least Squares Discriminant Analysis

Partial Least Squares

225 samples

91 predictor

3 classes: 'Mild', 'None', 'Severe'

Pre-processing: centered (91), scaled (91)

Resampling: Repeated Train/Test Splits Estimated (10 reps, 75%)


Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...

Resampling results across tuning parameters:

ncomp	Accuracy	Kappa
1	0.5035714	0.042641245
2	0.4839286	0.008444541
3	0.4589286	-0.022063946
4	0.4696429	0.005403703
5	0.4642857	-0.006616683
6	0.4607143	-0.008240888
7	0.4589286	-0.005506676
8	0.4357143	-0.038615293
9	0.4339286	-0.037265952
10	0.4464286	-0.014036123

Kappa was used to select the optimal model using the largest value.

The final value used for the model was ncomp = 1.



Confusion Matrix and Statistics

Prediction	Reference		
	Mild	None	Severe
Mild	26	17	6
None	3	4	0
Severe	0	0	0

Overall Statistics

Accuracy : 0.5357
95% CI : (0.3974, 0.6701)
No Information Rate : 0.5179
P-Value [Acc > NIR] : 0.4475

Kappa : 0.0714

McNemar's Test P-Value : NA

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.8966	0.19048	0.0000
Specificity	0.1481	0.91429	1.0000
Pos Pred Value	0.5306	0.57143	NaN
Neg Pred Value	0.5714	0.65306	0.8929
Prevalence	0.5179	0.37500	0.1071
Detection Rate	0.4643	0.07143	0.0000
Detection Prevalence	0.8750	0.12500	0.0000
Balanced Accuracy	0.5223	0.55238	0.5000

- Penalized Model

glmnet

225 samples

91 predictor

3 classes: 'Mild', 'None', 'Severe'

Pre-processing: centered (91), scaled (91)

Resampling: Repeated Train/Test Splits Estimated (10 reps, 75%)


Summary of sample sizes: 169, 169, 169, 169, 169, 169, ...

Resampling results across tuning parameters:

alpha	lambda	Accuracy	Kappa
0.0	0.0100	0.4517857	0.039628481
0.0	0.0575	0.4589286	0.009050681
0.0	0.1050	0.4714286	0.012617514
0.0	0.1525	0.4821429	0.024918133
0.0	0.2000	0.4785714	0.009809834
0.1	0.0100	0.4589286	0.044996563
0.1	0.0575	0.4535714	-0.013809616
0.1	0.1050	0.4714286	-0.005126396
0.1	0.1525	0.4803571	-0.002940312
0.1	0.2000	0.4767857	-0.021012823
0.2	0.0100	0.4642857	0.048835536
0.2	0.0575	0.4642857	-0.008623359
0.2	0.1050	0.4875000	0.011839575
0.2	0.1525	0.4839286	-0.013643231
0.2	0.2000	0.4875000	-0.019646290
0.4	0.0100	0.4482143	0.018221615
0.4	0.0575	0.4892857	0.022744838
0.4	0.1050	0.4946429	0.001897383
0.4	0.1525	0.4982143	-0.015268890
0.4	0.2000	0.5107143	-0.008420766
0.6	0.0100	0.4428571	-0.002084610
0.6	0.0575	0.4892857	0.006673787
0.6	0.1050	0.4982143	-0.015268890
0.6	0.1525	0.5196429	0.007308085
0.6	0.2000	0.5178571	0.000000000

Kappa was used to select the optimal model using the largest value.

The final values used for the model were alpha = 0.2 and lambda = 0.01.



Confusion Matrix and Statistics

	Reference		
Prediction	Mild	None	Severe
Mild	18	9	4
None	11	11	1
Severe	0	1	1

Overall Statistics

Accuracy : 0.5357
 95% CI : (0.3974, 0.6701)
 No Information Rate : 0.5179
 P-Value [Acc > NIR] : 0.4475

Kappa : 0.1642

Mcnemar's Test P-Value : 0.2407

Statistics by Class:

	Class: Mild	Class: None	Class: Severe
Sensitivity	0.6207	0.5238	0.16667
Specificity	0.5185	0.6571	0.98000
Pos Pred Value	0.5806	0.4783	0.50000
Neg Pred Value	0.5600	0.6970	0.90741
Prevalence	0.5179	0.3750	0.10714
Detection Rate	0.3214	0.1964	0.01786
Detection Prevalence	0.5536	0.4107	0.03571
Balanced Accuracy	0.5696	0.5905	0.57333

Model	Tuning parameter	Kappa
Linear Discriminant Analysis	N/A	0.1413
Partial Least Squares Discriminant Analysis	Components = 1	0.0714
Penalized	Alpha = 0.2, Lambda = 0.01	0.1642

- The penalized model with an alpha of 0.2 and a lambda of 0.01 was shown to be the highest performing model.

Problem 12.1d For the optimal model for the biological predictors, what are the top five important predictors?


- The top 5 most important biological predictors in the highest performing penalized model was shown to be Z167, Z38, Z42, Z100, and Z71.

glmnet variable importance


variables are sorted by maximum importance across the classes
only 20 most important variables shown (out of 91)

	Mild	None	Severe
Z167	100.00000	2.29	66.425
Z38	98.71615	61.65	5.785
Z42	0.00000	67.04	96.107
Z100	7.90491	43.57	82.759
Z71	39.34312	74.64	4.015
Z116	70.36695	35.27	3.810
Z36	0.00000	70.25	42.607
Z121	14.26399	66.12	20.576
Z59	28.98512	3.88	64.150
Z82	62.56418	22.62	8.657
Z145	3.12327	60.47	26.064
Z149	53.48731	60.28	0.000
Z128	4.75792	56.86	20.818
Z126	55.88285	12.47	12.124
Z15	0.05943	55.35	24.262
Z147	54.00999	0.00	28.040
Z132	51.60417	27.95	0.000
Z7	3.30901	16.71	51.304
Z176	40.24023	51.15	0.000
Z44	0.18084	50.58	19.889

Problem 12.3a Explore the data by visualizing the relationship between the predictors and the outcome. Are there important features of the predictor data themselves, such as between-predictor correlations or degenerate distributions?



- Only the number of voicemail messages was shown to be a low variance predictor, the histogram shows that the frequency is concentrated around zero. This will not add much to the predictive quality of the model so this predictor will be removed



- Some predictors do appear to be highly correlated in this dataset. The Correlation plot shows that:
 - Total evening minutes is associated with total eve charge
 - Total international minutes is associated with total international charge
 - Total night minutes is associated with total night charge
 - Total day minutes is associated with total day charge

Problem 12.3b What criteria should be used to evaluate the effectiveness of the models?

- The area under the ROC curve would be a good statistic to use to judge the models. The area under the ROC curve will maximize sensitivity and specificity, both of which we care about here.

Problem 12.3c Split the data into training set and test set using random splitting (80% and 20%). Fit models covered in this chapter to the training set and tune them via resampling. Which model has the best performance?

- Logistic Regression model

Generalized Linear Model

4001 samples

14 predictor

2 classes: 'yes', 'no'

Pre-processing: centered (11), scaled (11), ignore (3)

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results:

ROC	Sens	Spec
0.7569928	0.1582979	0.9797203

Confusion Matrix and Statistics

	Reference	
Prediction	yes	no
yes	18	19
no	123	839

Accuracy : 0.8579

95% CI : (0.8346, 0.8789)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.5584

Kappa : 0.1525

Mcnemar's Test P-Value : <2e-16

Sensitivity : 0.12766

Specificity : 0.97786

Pos Pred Value : 0.48649

Neg Pred Value : 0.87214

Prevalence : 0.14114

Detection Rate : 0.01802

Detection Prevalence : 0.03704

Balanced Accuracy : 0.55276

'Positive' Class : yes

- Penalized model

glmnet

4001 samples

14 predictor

2 classes: 'yes', 'no'

Pre-processing: centered (11), scaled (11), ignore (3)

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)


Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results across tuning parameters:

alpha	lambda	ROC	Sens	Spec
0.0	0.0100	0.7548930	0.0578723404	0.9897902
0.0	0.0575	0.7550470	0.0161702128	0.9984149
0.0	0.1050	0.7547534	0.0056737589	1.0000000
0.0	0.1525	0.7543266	0.0039716312	1.0000000
0.0	0.2000	0.7538994	0.0019858156	1.0000000
0.1	0.0100	0.7549515	0.0550354610	0.9904429
0.1	0.0575	0.7549349	0.0121985816	0.9991608
0.1	0.1050	0.7527726	0.0045390071	1.0000000
0.1	0.1525	0.7477664	0.0000000000	1.0000000
0.1	0.2000	0.7415533	0.0000000000	1.0000000
0.2	0.0100	0.7549617	0.0530496454	0.9910023
0.2	0.0575	0.7517992	0.0079432624	0.9994872
0.2	0.1050	0.7405512	0.0002836879	1.0000000
0.2	0.1525	0.7347817	0.0000000000	1.0000000
0.2	0.2000	0.7329399	0.0000000000	1.0000000
0.4	0.0100	0.7547868	0.0470921986	0.9924942
0.4	0.0575	0.7386677	0.0056737589	0.9998601
0.4	0.1050	0.7329749	0.0000000000	1.0000000
0.4	0.1525	0.7152967	0.0000000000	1.0000000
0.4	0.2000	0.5166475	0.0000000000	1.0000000
0.6	0.0100	0.7542859	0.0417021277	0.9932401
0.6	0.0575	0.7348348	0.0008510638	1.0000000
0.6	0.1050	0.7098031	0.0000000000	1.0000000
0.6	0.1525	0.5000000	0.0000000000	1.0000000
0.6	0.2000	0.5000000	0.0000000000	1.0000000

ROC was used to select the optimal model using the largest value.

The final values used for the model were alpha = 0 and lambda = 0.0575.



Confusion Matrix and Statistics

	Reference	
Prediction	yes	no
yes	2	2
no	139	856

Accuracy : 0.8589

95% CI : (0.8357, 0.8799)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.5224

Kappa : 0.02

Mcnemar's Test P-Value : <2e-16

Sensitivity : 0.014184

Specificity : 0.997669

Pos Pred Value : 0.500000

Neg Pred Value : 0.860302

Prevalence : 0.141141

Detection Rate : 0.002002

Detection Prevalence : 0.004004

Balanced Accuracy : 0.505927

'Positive' Class : yes

- Neural Net model

Neural Network

4001 samples

14 predictor

2 classes: 'yes', 'no'

Pre-processing: centered (11), scaled (11), spatial sign transformation (11), ignore (3)

Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)


Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results across tuning parameters:

size	decay	ROC	Sens	Spec
1	0.0	0.6609511	0.02666667	0.9977622
1	0.1	0.7306211	0.14212766	0.9823310
1	0.3	0.7452853	0.09900709	0.9846620
1	0.5	0.7469900	0.08482270	0.9875991
1	1.0	0.7535754	0.05134752	0.9937995
2	0.0	0.6852687	0.11659574	0.9740793
2	0.1	0.7734282	0.30269504	0.9671795
2	0.3	0.7856192	0.29106383	0.9733800
2	0.5	0.7893997	0.25248227	0.9806061
2	1.0	0.8166179	0.24822695	0.9888578
3	0.0	0.6473379	0.19375887	0.9517949
3	0.1	0.7899284	0.37049645	0.9627972
3	0.3	0.8146810	0.37787234	0.9717949
3	0.5	0.8166982	0.34609929	0.9773427
3	1.0	0.8191139	0.28028369	0.9867133

ROC was used to select the optimal model using the largest value.

The final values used for the model were size = 3 and decay = 1.



Confusion Matrix and Statistics

Reference
Prediction yes no
yes 44 10
no 97 848

Accuracy : 0.8929
95% CI : (0.872, 0.9114)
No Information Rate : 0.8589
P-Value [Acc > NIR] : 0.0008284

Kappa : 0.4048

Mcnemar's Test P-Value : < 2.2e-16

Sensitivity : 0.31206
Specificity : 0.98834
Pos Pred Value : 0.81481
Neg Pred Value : 0.89735
Prevalence : 0.14114
Detection Rate : 0.04404
Detection Prevalence : 0.05405
Balanced Accuracy : 0.65020

'Positive' Class : yes

- Mixed Discriminant Analysis

Mixture Discriminant Analysis

4001 samples

14 predictor

2 classes: 'yes', 'no'

No pre-processing


Resampling: Repeated Train/Test Splits Estimated (25 reps, 75%)

Summary of sample sizes: 3002, 3002, 3002, 3002, 3002, 3002, ...

Resampling results across tuning parameters:

subclasses	ROC	Sens	Spec
1	0.7623251	0.1520567	0.9778555
2	0.8272049	0.4252482	0.9765967
3	0.8188054	0.4184397	0.9686247
4	0.8165392	0.4592908	0.9569697

ROC was used to select the optimal model using the largest value.
The final value used for the model was subclasses = 2.



Confusion Matrix and Statistics

Reference

Prediction yes no

yes 43 16

no 98 842

Accuracy : 0.8859

95% CI : (0.8645, 0.9049)

No Information Rate : 0.8589

P-Value [Acc > NIR] : 0.006838

Kappa : 0.3782

Mcnemar's Test P-Value : 3.291e-14

Sensitivity : 0.30496

Specificity : 0.98135

Pos Pred Value : 0.72881

Neg Pred Value : 0.89574

Prevalence : 0.14114

Detection Rate : 0.04304

Detection Prevalence : 0.05906

Balanced Accuracy : 0.64316

'Positive' Class : yes

Model	Tuning Parameters	ROC
Logistic Regression	N/A	0.736
Penalized model	Alpha = 0, lambda = 0.0575	0.7647
Neural Net	Size = 3, decay = 1	0.7984
Mixed Discriminant Analysis	Subclasses = 2	0.7913

- The Neural Net model with size = 3 and decay = 1 performed the best of any of the models tested here as it had the highest area under the ROC curve.

R Code

#Assignment Week 6

#Cody Rorick

library(caret)

library(AppliedPredictiveModeling)

data(hepatic)

#Problem 12.1c Pre-process the data, split the data into a training and a testing set, and build models described

#in this chapter for the biological predictors. Using each model to predict on the testing set, which

#model has the best predictive ability for the biological predictors and what is the optimal

#performance?

low.variability.columns <- nearZeroVar(bio)

bio <- bio[,-low.variability.columns]

correlations <- cor(bio)

library(corrplot)

corrplot(correlations, order = "hclust")

highly.correlated.columns <- findCorrelation(correlations, cutoff = .85)

bio <- bio[,-highly.correlated.columns]

set.seed(100)

trainRows <- createDataPartition(injury, p = .80, list = FALSE)

trainResp <- injury[trainRows]

trainPred <- bio[trainRows,]

testResp <- injury[-trainRows]

testPred <- bio[-trainRows,]

Linear Discriminant Analysis

ctrl <- trainControl(method = "LGOCV", number = 10, classProbs = TRUE, savePredictions = TRUE)

LDAFull <- train(trainPred, y = trainResp, method = "lda", preProc = c("center","scale"), metric = "Kappa", trControl = ctrl)

LDAFull

confusionMatrix(data = predict(LDAFull, testPred), reference = testResp) #average over 10*.25*1000 observations

Partial Least Squares Discriminant Analysis

set.seed(100)

ctrl <- trainControl(method = "LGOCV", number = 10, classProbs = TRUE, savePredictions = TRUE)

plsFit <- train(x = trainPred, y = trainResp, method = "pls", tuneGrid = expand.grid(.ncomp = 1:10), preProc = c("center","scale"), metric = "Kappa", trControl = ctrl)

plsFit


```
plot(plsFit)
confusionMatrix(data = predict(plsFit, testPred), reference = testResp)
```

```
####Penalized glmnet model
```

```
library(glmnet)
ctrl <- trainControl(method = "LGOCV", number = 10, classProbs = TRUE, savePredictions =
TRUE)
glmnetGrid <- expand.grid(.alpha = c(0, .1, .2, .4, .6), .lambda = seq(.01, .2, length = 5))
set.seed(100)
glmnetTuned <- train(trainPred, y = trainResp, method = "glmnet", tuneGrid = glmnetGrid, preProc =
c("center", "scale"), metric = "Kappa", trControl = ctrl)
glmnetTuned
plot(glmnetTuned)
confusionMatrix(data = predict(glmnetTuned, testPred), reference = testResp)
```

```
#Problem 12.1d For the optimal model for the biological predictors, what are the top five
important predictors?
varImp(glmnetTuned)
```

```
#Problem 12.3a Explore the data by visualizing the relationship between the predictors and the
outcome. Are
```

```
#there important features of the predictor data themselves, such as between-predictor
correlations
```

```
#or degenerate distributions?
```

```
library(modeldata)
data("mlc_churn")
low.variability.columns <- nearZeroVar(mlc_churn)
hist(unlist(mlc_churn[,6]), main = 'Vmail Messages', xlab = '# Vmail Messages')
mlc_churn <- mlc_churn[, -low.variability.columns]
correlations <- cor(mlc_churn[, -c(1,3,4,5,19)])
corrplot(correlations, order = "hclust")
highly.correlated.columns <- findCorrelation(correlations, cutoff = .85)
mlc_churn <- mlc_churn[, -highly.correlated.columns]
```

```
#Problem 12.3c Split the data into training set and test set using random splitting (80% and
20%). Fit models
```

```
#covered in this chapter to the training set and tune them via resampling. Which model has the
best performance?
```

```
set.seed(100)
trainRows <- createDataPartition(mlc_churn$churn, p = .80, list = FALSE)
trainResp <- mlc_churn$churn[trainRows]
trainPred <- mlc_churn[trainRows, -c(15)]
testResp <- mlc_churn$churn[-trainRows]
```

```
testPred <- mlc_churn[-trainRows,-c(15)]
```

Logistic Regression

```
set.seed(100)
ctrl <- trainControl(method = "LGOCV", summaryFunction = twoClassSummary, classProbs =
TRUE, savePredictions = TRUE)
lrFull <- train(trainPred, y = trainResp, method = "glm",preProc = c("center", "scale"), metric =
"ROC", trControl = ctrl)
lrFull
confusionMatrix(data = predict(lrFull, testPred), reference = testResp)
library(pROC)
rocCurve <- roc(response = testResp, predictor = predict(lrFull, testPred, type = 'prob')[,2])
auc(rocCurve)
```

Penalized glmnet model

```
library(glmnet)
ctrl <- trainControl(method = "LGOCV",summaryFunction = twoClassSummary,classProbs =
TRUE,savePredictions = TRUE)
glmnetGrid <- expand.grid(.alpha = c(0, .1, .2, .4, .6),.lambda = seq(.01, .2, length = 5))
set.seed(100)
glmnetTuned <- train(trainPred, y = trainResp, method = "glmnet",tuneGrid = glmnetGrid,preProc =
c("center", "scale"),metric = "ROC",trControl = ctrl)
glmnetTuned
plot(glmnetTuned)
confusionMatrix(data = predict(glmnetTuned, testPred), reference = testResp)
rocCurve <- roc(response = testResp, predictor = predict(glmnetTuned, testPred, type =
'prob')[,1])
auc(rocCurve)
```

Neural Net Model

```
nnetGrid <- expand.grid(.size = 1:3, .decay = c(0, .1, .3, .5, 1))
maxSize <- max(nnetGrid$.size)
numWts <- (15 * (14 + 1) + (15+1)*2) ## 14 is the number of predictors; 2 is the number of
classes; 15 is size*decay
ctrl <- trainControl(method = 'LGOCV', summaryFunction = twoClassSummary,classProbs =
TRUE)
nnetFit <- train(x = trainPred, y = trainResp,method = "nnet",metric = "ROC",preProc =
c("center", "scale", "spatialSign"),tuneGrid = nnetGrid,trace = FALSE,maxit = 2000,MaxNWts =
numWts,trControl = ctrl)
nnetFit
plot(nnetFit)
confusionMatrix(data = predict(nnetFit, testPred), reference = testResp)
rocCurve <- roc(response = testResp, predictor = predict(nnetFit, testPred, type = 'prob')[,1])
auc(rocCurve)
```

```
### Mixed Discriminant Analysis
library(mda)
ctrl <- trainControl(method = "LGOCV", summaryFunction = twoClassSummary, classProbs =
TRUE, savePredictions = TRUE)
set.seed(100)
mdaFit <- train(x = trainPred, y = trainResp, method = "mda", metric = "ROC", tuneGrid =
expand.grid(.subclasses = 1:4), trControl = ctrl)
mdaFit
plot(mdaFit)
confusionMatrix(data = predict(mdaFit, testPred), reference = testResp)
rocCurve <- roc(response = testResp, predictor = predict(mdaFit, testPred, type = 'prob')[,1])
auc(rocCurve)
```